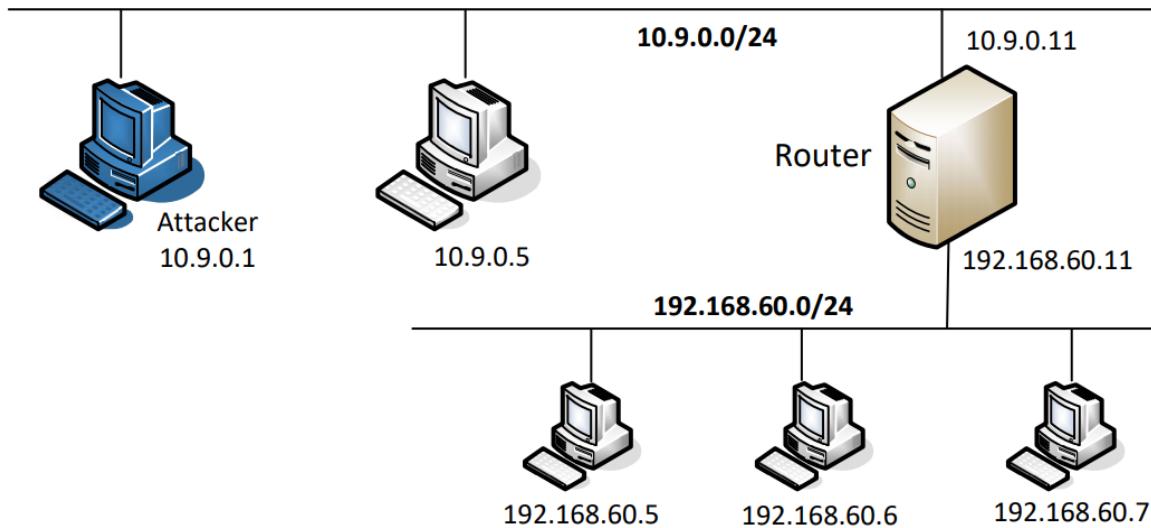


# lab6 Firewall Exploration Lab

51778204 陈盈

## Environment Setup Using Containers



查看各主机的哈希值。

```
d [07/26/21] seed@VM:~/.../Labsetup$ dockps
1bc65e6feaa8  hostA-10.9.0.5
1d6463215ecb  host1-192.168.60.5
8068735d0b54  host2-192.168.60.6
5b0a2c8e5bd9  host3-192.168.60.7
6c05b2e686e8  seed-router
```

## Task 1: Implementing a Simple Firewall

### Task 1.A: Implement a Simple Kernel Module

对kernel\_modules进行编译。

```
[07/26/21] seed@VM:~/kernel_module$ make
make -C /lib/modules/5.4.0-54-generic/build M=/home/seed/kernel_module modules
make[1]: Entering directory '/usr/src/linux-headers-5.4.0-54-generic'
  Building modules, stage 2.
    MODPOST 1 modules
WARNING: modpost: missing MODULE_LICENSE() in /home/seed/kernel_module/hello.o
see include/linux/module.h for more information
make[1]: Leaving directory '/usr/src/linux-headers-5.4.0-54-generic'
```

测试以下命令。

```
[07/26/21] seed@VM:~/kernel_module$ sudo insmod hello.ko
[07/26/21] seed@VM:~/kernel_module$ lsmod | grep hello
hello          16384  0
[07/26/21] seed@VM:~/kernel_module$ dmesg | grep World
[ 1220.434008] Hello World!
[ 1363.966704] Bye-bye World!.
```

## Task 1.B: Implement a Simple Firewall Using Netfilter

1

和task1.A一样，对文件进行编译。

```
[07/26/21]seed@VM:~/packet_filter$ make
make -C /lib/modules/5.4.0-54-generic/build M=/home/seed/packet_filter modules
make[1]: Entering directory '/usr/src/linux-headers-5.4.0-54-generic'
  CC [M]  /home/seed/packet_filter/seedFilter.o
  Building modules, stage 2.
MODPOST 1 modules
  CC [M]  /home/seed/packet_filter/seedFilter.mod.o
  LD [M]  /home/seed/packet_filter/seedFilter.ko
make[1]: Leaving directory '/usr/src/linux-headers-5.4.0-54-generic'
```

加载内核前，可以看到 dig @8.8.8.8 www.example 命令可以得到响应。

```
[07/26/21]seed@VM:~/packet_filter$ dig @8.8.8.8 www.example.com
; <>> DiG 9.16.1-Ubuntu <>> @8.8.8.8 www.example.com
; (1 server found)
; global options: +cmd
; Got answer:
; >>>HEADER<<- opcode: QUERY, status: NOERROR, id: 55651
; flags: qr rd ra ad; QUERY: 1, ANSWER: 1, AUTHORITY: 0, ADDITIONAL: 1

; OPT PSEUDOSECTION:
; EDNS: version: 0, flags:; udp: 512
; QUESTION SECTION:
;www.example.com.           IN      A

; ANSWER SECTION:
www.example.com.        20148    IN      A      93.184.216.34

; Query time: 84 msec
; SERVER: 8.8.8.8#53(8.8.8.8)
; WHEN: Mon Jul 26 11:50:18 EDT 2021
; MSG SIZE  rcvd: 60
```

加载内核后，防火墙生效，dig @8.8.8.8 www.example 命令得不到响应。

```
[07/26/21]seed@VM:~/packet_filter$ sudo insmod seedFilter.ko
[07/26/21]seed@VM:~/packet_filter$ lsmod | grep seedFilter
seedFilter            16384   0
[07/26/21]seed@VM:~/packet_filter$ dig @8.8.8.8 www.example.com
; <>> DiG 9.16.1-Ubuntu <>> @8.8.8.8 www.example.com
; (1 server found)
; global options: +cmd
; connection timed out; no servers could be reached
```

完成任务后移除内核。

```
[07/26/21]seed@VM:~/packet_filter$ sudo rmmod seedFilter
[07/26/21]seed@VM:~/packet_filter$ lsmod | grep seedFilter
```

2

进行 dig @8.8.8.8 [www.example.com](http://www.example.com) 操作后，可使用 sudo dmesg -c 查看信息，每次测试后，需要运行 sudo rmmod seedFilter 从内核中移除模块。

未加载内核时dig @8.8.8.8 [www.example.com](http://www.example.com)。

```
[07/26/21]seed@VM:~/packet_filter$ sudo dmesg -c
[ 2143.859230] *** LOCAL_OUT
[ 2143.859233] 172.17.0.1 --> 224.0.0.251 (UDP)
[ 2194.107175] *** LOCAL_OUT
[ 2194.107178] 192.168.43.199 --> 224.0.0.251 (UDP)
[ 2196.193443] *** LOCAL_OUT
[ 2196.193448] 192.168.43.199 --> 192.168.43.1 (UDP)
[ 2196.252647] *** LOCAL_OUT
[ 2196.252655] 127.0.0.1 --> 127.0.0.53 (UDP)
[ 2196.253127] *** LOCAL_OUT
[ 2196.253129] 192.168.43.199 --> 192.168.43.1 (UDP)
[ 2196.258034] *** LOCAL_OUT
[ 2196.258039] 127.0.0.53 --> 127.0.0.1 (UDP)
[ 2200.429576] The filters are being removed.
```

修改seedFilter.c文件，代码如下。

```

#include <linux/kernel.h>
#include <linux/module.h>
#include <linux/netfilter.h>
#include <linux/netfilter_ipv4.h>
#include <linux/ip.h>
#include <linux/tcp.h>
#include <linux/udp.h>
#include <linux/if_ether.h>
#include <linux/inet.h>

static struct nf_hook_ops hook1, hook2, hook3, hook4, hook5;

unsigned int blockUDP(void *priv, struct sk_buff *skb,
                      const struct nf_hook_state *state)
{
    struct iphdr *iph;
    struct udphdr *udph;

    u16 port = 53;
    char ip[16] = "8.8.8.8";
    u32 ip_addr;

    if (!skb) return NF_ACCEPT;

    iph = ip_hdr(skb);
    // Convert the IPv4 address from dotted decimal to 32-bit binary
    in4_pton(ip, -1, (u8 *)&ip_addr, '\0', NULL);

    if (iph->protocol == IPPROTO_UDP) {
        udph = udp_hdr(skb);
        if (iph->daddr == ip_addr && ntohs(udph->dest) == port){
            printk(KERN_WARNING "*** Dropping %pI4 (UDP), port %d\n",
                  &(iph->daddr), port);
            return NF_DROP;
        }
    }
    return NF_ACCEPT;
}

unsigned int printInfo(void *priv, struct sk_buff *skb,
                      const struct nf_hook_state *state)
{
    struct iphdr *iph;
    char *hook;
    char *protocol;

    switch (state->hook){
        case NF_INET_LOCAL_IN:      hook = "LOCAL_IN";      break;
        case NF_INET_LOCAL_OUT:     hook = "LOCAL_OUT";     break;
        case NF_INET_PRE_ROUTING:   hook = "PRE_ROUTING";   break;
        case NF_INET_POST_ROUTING:  hook = "POST_ROUTING";  break;
        case NF_INET_FORWARD:       hook = "FORWARD";       break;
        default:                   hook = "IMPOSSIBLE";    break;
    }
    printk(KERN_INFO "*** %s\n", hook); // Print out the hook info
}

```

```

iph = ip_hdr(skb);
switch (iph->protocol){
    case IPPROTO_UDP: protocol = "UDP"; break;
    case IPPROTO_TCP: protocol = "TCP"; break;
    case IPPROTO_ICMP: protocol = "ICMP"; break;
    default:           protocol = "OTHER"; break;
}

// Print out the IP addresses and protocol
printf(KERN_INFO "%pI4 --> %pI4 (%s)\n",
       &(iph->saddr), &(iph->daddr), protocol);

return NF_ACCEPT;
}

int registerFilter(void) {
printf(KERN_INFO "Registering filters.\n");
// Hook 1
hook1.hook = printInfo;
hook1.hooknum = NF_INET_LOCAL_IN;
hook1(pf = PF_INET;
hook1.priority = NF_IP_PRI_FIRST;
nf_register_net_hook(&init_net, &hook1);
// Hook 2
hook2.hook = printInfo;
hook2.hooknum = NF_INET_PRE_ROUTING;
hook2(pf = PF_INET;
hook2.priority = NF_IP_PRI_FIRST;
nf_register_net_hook(&init_net, &hook2);
// Hook 3
hook3.hook = printInfo;
hook3.hooknum = NF_INET_FORWARD;
hook3(pf = PF_INET;
hook3.priority = NF_IP_PRI_FIRST;
nf_register_net_hook(&init_net, &hook3);
// Hook 4
hook4.hook = printInfo;
hook4.hooknum = NF_INET_LOCAL_OUT;
hook4(pf = PF_INET;
hook4.priority = NF_IP_PRI_FIRST;
nf_register_net_hook(&init_net, &hook4);
// Hook 5
hook5.hook = printInfo;
hook5.hooknum = NF_INET_POST_ROUTING;
hook5(pf = PF_INET;
hook5.priority = NF_IP_PRI_FIRST;
nf_register_net_hook(&init_net, &hook5);
return 0;
}
void removeFilter(void) {
printf(KERN_INFO "The filters are being removed.\n");
nf_unregister_net_hook(&init_net, &hook1);
nf_unregister_net_hook(&init_net, &hook2);
nf_unregister_net_hook(&init_net, &hook3);
nf_unregister_net_hook(&init_net, &hook4);
}

```

```

nf_unregister_net_hook(&init_net, &hook5);
}
module_init(registerFilter);
module_exit(removeFilter);
MODULE_LICENSE("GPL");

```

编译文件并装载内核。

```

[07/26/21]seed@VM:~/packet_filter$ make
make -C /lib/modules/5.4.0-54-generic/build M=/home/seed/packet_filter modules
make[1]: Entering directory '/usr/src/linux-headers-5.4.0-54-generic'
  CC [M]  /home/seed/packet_filter/seedFilter.o
  Building modules, stage 2.
MODPOST 1 modules
  CC [M]  /home/seed/packet_filter/seedFilter.mod.o
  LD [M]  /home/seed/packet_filter/seedFilter.ko
make[1]: Leaving directory '/usr/src/linux-headers-5.4.0-54-generic'
[07/26/21]seed@VM:~/packet_filter$ sudo insmod seedFilter.ko
[07/26/21]seed@VM:~/packet_filter$ lsmod | grep seedFilter
seedFilter           16384  0

```

在用户主机上ping攻击者主机，得到结果如下，可知能够连接。

```

root@1bc65e6feaa8:/# ping 10.9.0.1
PING 10.9.0.1 (10.9.0.1) 56(84) bytes of data.
64 bytes from 10.9.0.1: icmp_seq=1 ttl=64 time=1.13 ms
64 bytes from 10.9.0.1: icmp_seq=2 ttl=64 time=0.091 ms
64 bytes from 10.9.0.1: icmp_seq=3 ttl=64 time=0.058 ms
64 bytes from 10.9.0.1: icmp_seq=4 ttl=64 time=0.259 ms
64 bytes from 10.9.0.1: icmp_seq=5 ttl=64 time=0.279 ms
64 bytes from 10.9.0.1: icmp_seq=6 ttl=64 time=0.073 ms
64 bytes from 10.9.0.1: icmp_seq=7 ttl=64 time=0.070 ms
64 bytes from 10.9.0.1: icmp_seq=8 ttl=64 time=0.498 ms
64 bytes from 10.9.0.1: icmp_seq=9 ttl=64 time=0.057 ms
64 bytes from 10.9.0.1: icmp_seq=10 ttl=64 time=0.072 ms
64 bytes from 10.9.0.1: icmp_seq=11 ttl=64 time=0.068 ms
64 bytes from 10.9.0.1: icmp_seq=12 ttl=64 time=0.379 ms

```

```

[07/26/21]seed@VM:~/packet_filter$ sudo dmesg -c
[ 4455.004360] e1000: ens33 NIC Link is Down
[ 4455.419151] br-670f453a63a6: port 2(vethfa2fce3) entered disabled state
[ 4455.419207] br-670f453a63a6: port 1(vethb3f05a6) entered disabled state
[ 4455.448386] br-83e44f406fe2: port 4(vethbdb9837) entered disabled state
[ 4455.448438] br-83e44f406fe2: port 3(veth5ddd0b9) entered disabled state
[ 4455.448445] br-83e44f406fe2: port 2(veth159cb26) entered disabled state
[ 4455.448450] br-83e44f406fe2: port 1(veth23435e2) entered disabled state
[ 4666.884812] e1000: ens33 NIC Link is Up 1000 Mbps Full Duplex, Flow Control:
None
[ 4803.779857] Registering filters.
[ 4929.217851] veth8202d13: renamed from eth0
[ 4929.244879] veth266b61e: renamed from eth0
[ 4929.259105] veth66b1342: renamed from eth0
[ 4929.271435] veth6472b85: renamed from eth0
[ 4929.304390] veth399a99c: renamed from eth0
[ 4929.384873] device veth23435e2 left promiscuous mode
[ 4929.384878] br-83e44f406fe2: port 1(veth23435e2) entered disabled state
[ 4929.425656] device vethb3f05a6 left promiscuous mode
[ 4929.425670] br-670f453a63a6: port 1(vethb3f05a6) entered disabled state
[ 4929.536765] device veth159cb26 left promiscuous mode
[ 4929.536771] br-83e44f406fe2: port 2(veth159cb26) entered disabled state
[ 4929.606439] device vethfa2fce3 left promiscuous mode
[ 4929.606444] br-670f453a63a6: port 2(vethfa2fce3) entered disabled state
[ 4929.660145] device veth5ddd0b9 left promiscuous mode
[ 4929.660150] br-83e44f406fe2: port 3(veth5ddd0b9) entered disabled state
[ 4929.834385] veth6be52e6: renamed from eth1
[ 4929.921655] device vethbdb9837 left promiscuous mode
[ 4929.921660] br-83e44f406fe2: port 4(vethbdb9837) entered disabled state

```

```
[ 5021.665526] *** PRE_ROUTING
[ 5021.665528]      10.9.0.5 --> 10.9.0.1 (ICMP)
[ 5021.665538] *** LOCAL_IN
[ 5021.665539]      10.9.0.5 --> 10.9.0.1 (ICMP)
[ 5021.665621] *** LOCAL_OUT
[ 5021.665623]      10.9.0.1 --> 10.9.0.5 (ICMP)
[ 5021.665630] *** POST_ROUTING
[ 5021.665631]      10.9.0.1 --> 10.9.0.5 (ICMP)
[ 5022.666623] *** PRE_ROUTING
[ 5022.666627]      10.9.0.5 --> 10.9.0.1 (ICMP)
[ 5022.666636] *** PRE_ROUTING
[ 5022.666637]      10.9.0.5 --> 10.9.0.1 (ICMP)
[ 5022.666646] *** LOCAL_IN
[ 5022.666646]      10.9.0.5 --> 10.9.0.1 (ICMP)
[ 5022.666653] *** LOCAL_OUT
[ 5022.666653]      10.9.0.1 --> 10.9.0.5 (ICMP)
[ 5022.666655] *** POST_ROUTING
[ 5022.666656]      10.9.0.1 --> 10.9.0.5 (ICMP)
[ 5023.676416] *** PRE_ROUTING
[ 5023.676418]      10.9.0.5 --> 10.9.0.1 (ICMP)
[ 5023.676423] *** PRE_ROUTING
[ 5023.676423]      10.9.0.5 --> 10.9.0.1 (ICMP)
[ 5023.676428] *** LOCAL_IN
[ 5023.676429]      10.9.0.5 --> 10.9.0.1 (ICMP)
[ 5023.676434] *** LOCAL_OUT
```

```
[07/26/21]seed@VM:~/packet_filter$ sudo rmmod seedFilter
[07/26/21]seed@VM:~/packet_filter$ lsmod | grep seedFilter
```

ping外部主机

```
root@1bc65e61eaa8:/# ping 192.168.60.5
PING 192.168.60.5 (192.168.60.5) 56(84) bytes of data.
64 bytes from 192.168.60.5: icmp_seq=1 ttl=63 time=0.114 ms
64 bytes from 192.168.60.5: icmp_seq=2 ttl=63 time=0.452 ms
64 bytes from 192.168.60.5: icmp_seq=3 ttl=63 time=0.120 ms
64 bytes from 192.168.60.5: icmp_seq=4 ttl=63 time=0.087 ms
64 bytes from 192.168.60.5: icmp_seq=5 ttl=63 time=0.134 ms
64 bytes from 192.168.60.5: icmp_seq=6 ttl=63 time=0.275 ms
64 bytes from 192.168.60.5: icmp_seq=7 ttl=63 time=0.711 ms
64 bytes from 192.168.60.5: icmp_seq=8 ttl=63 time=0.102 ms
64 bytes from 192.168.60.5: icmp_seq=9 ttl=63 time=0.197 ms
64 bytes from 192.168.60.5: icmp_seq=10 ttl=63 time=0.287 ms
64 bytes from 192.168.60.5: icmp_seq=11 ttl=63 time=1.04 ms
64 bytes from 192.168.60.5: icmp_seq=12 ttl=63 time=0.465 ms
64 bytes from 192.168.60.5: icmp_seq=13 ttl=63 time=0.687 ms
64 bytes from 192.168.60.5: icmp_seq=14 ttl=63 time=0.101 ms
64 bytes from 192.168.60.5: icmp_seq=15 ttl=63 time=0.115 ms
64 bytes from 192.168.60.5: icmp_seq=16 ttl=63 time=0.080 ms
64 bytes from 192.168.60.5: icmp_seq=17 ttl=63 time=0.146 ms
```

```
[07/26/21] seed@VM:~/packet_filter$ sudo dmesg -c
[ 5067.837763] *** LOCAL_OUT
[ 5067.837765] 127.0.0.1 --> 127.0.0.53 (UDP)
[ 5067.837774] *** POST_ROUTING
[ 5067.837775] 127.0.0.1 --> 127.0.0.53 (UDP)
[ 5067.837789] *** PRE_ROUTING
[ 5067.837790] 127.0.0.1 --> 127.0.0.53 (UDP)
[ 5067.837791] *** LOCAL_IN
[ 5067.837791] 127.0.0.1 --> 127.0.0.53 (UDP)
[ 5067.838010] *** LOCAL_OUT
[ 5067.838012] 127.0.0.53 --> 127.0.0.1 (UDP)
[ 5067.838014] *** POST_ROUTING
[ 5067.838015] 127.0.0.53 --> 127.0.0.1 (UDP)
[ 5067.838019] *** PRE_ROUTING
[ 5067.838019] 127.0.0.53 --> 127.0.0.1 (UDP)
[ 5067.838020] *** LOCAL_IN
[ 5067.838020] 127.0.0.53 --> 127.0.0.1 (UDP)
[ 5067.838968] *** LOCAL_OUT
[ 5067.838969] 127.0.0.1 --> 127.0.0.53 (UDP)
[ 5067.838976] *** POST_ROUTING
[ 5067.838980] 127.0.0.1 --> 127.0.0.53 (UDP)
[ 5067.838987] *** PRE_ROUTING
[ 5067.838988] 127.0.0.1 --> 127.0.0.53 (UDP)
[ 5067.838989] *** LOCAL_IN
[ 5067.838989] 127.0.0.1 --> 127.0.0.53 (UDP)
[ 5067.839159] *** LOCAL_OUT
[ 5067.839160] 127.0.0.53 --> 127.0.0.1 (UDP)
[ 5067.839163] *** POST_ROUTING
[ 5067.839163] 127.0.0.53 --> 127.0.0.1 (UDP)
```

数据报从进入系统，进行 IP 校验以后，首先经过第一个 HOOK 函数 NF\_INET\_PRE\_ROUTING 进行处理，然后就进入路由代码，其决定该数据报是需要转发还是发给本机的。若该数据报应该被转发则它被 NF\_INET\_FORWARD 处理。

挂载 NF\_INET\_LOCAL\_OUT 时，本机产生的数据包将会第一个到达此 HOOK，数据经过 HOOK 函数 NF\_INET\_LOCAL\_OUT 处理后，进行路由选择处理，然后经过 NF\_INET\_POST\_ROUTING 处理后发送出去。

经过转发的数据报经过最后一个 HOOK 函数 NF\_INET\_POST\_ROUTING 处理以后，再传输到网络上。

### 3

修改seedFilter.c文件，代码如下：

```
#include <linux/kernel.h>
#include <linux/module.h>
#include <linux/netfilter.h>
#include <linux/netfilter_ipv4.h>
#include <linux/ip.h>
#include <linux/tcp.h>

static struct nf_hook_ops telnetFilterHook;

unsigned int telnetFilter(void *priv, struct sk_buff *skb, const struct nf_hook_state *state){

    struct iphdr *iph;
    struct tcphdr *tcpiph;
    iph = ip_hdr(skb);
    tcpiph = (void *)iph+iph->ihl*4;

    if((iph->protocol == IPPROTO_TCP && (tcpiph->dest == htons(23))
    || tcpiph->dest== htons(22))
```

```

        || tcph->dest == htons(21)))
        || (iph->protocol == IPPROTO_ICMP &&(((unsigned char *)&iph-
>daddr)[0]==10 &&
            ((unsigned char *)&iph->daddr)[1]==9
            && ((unsigned char *)&iph->daddr)[2]==0 && ((unsigned char
*)&iph->daddr)[3]==1)
            || (((unsigned char *)&iph->daddr)[0]==10 && ((unsigned char
*)&iph->daddr)[1]==9
            && ((unsigned char *)&iph->daddr)[2]==0 && ((unsigned char
*)&iph->daddr)[3]==1)))){
            printk(KERN_INFO "Dropping telent packdt to %d.%d.%d.%d\n",
            ((unsigned char *)&iph->daddr)[0],
            ((unsigned char *)&iph->daddr)[1],
            ((unsigned char *)&iph->daddr)[2],
            ((unsigned char *)&iph->daddr)[3]);
            return NF_DROP;
        }else{
            return NF_ACCEPT;
        }
    }
}

void removeFilter(void){
    printk(KERN_INFO "Telnet filter has been removed.\n");
    nf_unregister_net_hook(&init_net,&telnetFilterHook);
}

int setUpFilter(void){

    telnetFilterHook.hook = telnetFilter;
    telnetFilterHook.hooknum = NF_INET_PRE_ROUTING;
    telnetFilterHook(pf = PF_INET;
    telnetFilterHook.priority = NF_IP_PRI_FILTER;

    if(nf_register_net_hook(&init_net,&telnetFilterHook)!=0){
        printk(KERN_WARNING "register Telnet filter hook error!\n");
        goto err;
    }
    printk(KERN_INFO "Registering a Telnet filter");
    return 0;
}

err:
    removeFilter();
    return -1;
}

module_init(setUpFilter);
module_exit(removeFilter);

MODULE_LICENSE("GPL");

```

编译文件并加载内核。

```

[07/26/21]seed@VM:~/packet_filter$ make
make -C /lib/modules/5.4.0-54-generic/build M=/home/seed/packet_filter modules
make[1]: Entering directory '/usr/src/linux-headers-5.4.0-54-generic'
  CC [M]  /home/seed/packet_filter/seedFilter.o
Building modules, stage 2.
MODPOST 1 modules
  CC [M]  /home/seed/packet_filter/seedFilter.mod.o
  LD [M]  /home/seed/packet_filter/seedFilter.ko
make[1]: Leaving directory '/usr/src/linux-headers-5.4.0-54-generic'
[07/26/21]seed@VM:~/packet_filter$ sudo insmod seedFilter.ko
[07/26/21]seed@VM:~/packet_filter$ lsmod | grep seedFilter
seedFilter           16384  0

```

在主机A(10.9.0.5)分别进行 ping 10.9.0.1 和 telnet 10.9.0.1。

```
root@1bc65e6feaa8:/# ping 10.9.0.1
PING 10.9.0.1 (10.9.0.1) 56(84) bytes of data.
^C
--- 10.9.0.1 ping statistics ---
4 packets transmitted, 0 received, 100% packet loss, time 3071ms
```

```
root@1bc65e6feaa8:/# telnet 10.9.0.1
Trying 10.9.0.1...
^C
```

在本机上查看内核缓存。

```
[ 5708.670068] The filters are being removed.
[ 5726.005624] Registering a Telnet filter
[ 5761.403980] Dropping telent packdt to 10.9.0.1
[ 5762.428247] Dropping telent packdt to 10.9.0.1
[ 5763.453870] Dropping telent packdt to 10.9.0.1
[ 5764.475027] Dropping telent packdt to 10.9.0.1
[ 5804.006753] Dropping telent packdt to 10.9.0.1
[ 5805.017561] Dropping telent packdt to 10.9.0.1
[ 5807.035009] Dropping telent packdt to 10.9.0.1
```

## Task 2: Experimenting with Stateless Firewall Rules

每个任务前都要清理 table。

### Task 2.A: Protecting the Router

按要求在router上输入以下命令。

```
root@6c05b2e686e8:/# iptables -A INPUT -p icmp --icmp-type echo-reply -j ACCEPT
root@6c05b2e686e8:/# iptables -A OUTPUT -p icmp --icmp-type echo-request -j ACCEPT
root@6c05b2e686e8:/# iptables -P OUTPUT DROP
root@6c05b2e686e8:/# iptables -P INPUT DROP
```

在主机A (10.9.0.5) 上, ping 10.9.0.11 和 telnet 10.9.0.11 都不通。

```
root@1bc65e6feaa8:/# ping 10.9.0.11
PING 10.9.0.11 (10.9.0.11) 56(84) bytes of data.
^C
--- 10.9.0.11 ping statistics ---
4 packets transmitted, 0 received, 100% packet loss, time 3052ms

root@1bc65e6feaa8:/# telnet 10.9.0.11
Trying 10.9.0.11...
^C
```

将命令换成如下顺序。

```
iptables -A INPUT -p icmp --icmp-type echo-request -j ACCEPT
iptables -A OUTPUT -p icmp --icmp-type echo-reply -j ACCEPT
iptables -P INPUT DROP
iptables -P OUTPUT DROP
```

在主机A (10.9.0.5) 上, ping 10.9.0.11通, 但 telnet 10.9.0.11 不通。

```
PING 10.9.0.11 (10.9.0.11) 56(84) bytes of data.
64 bytes from 10.9.0.11: icmp_seq=1 ttl=64 time=0.134 ms
64 bytes from 10.9.0.11: icmp_seq=2 ttl=64 time=0.058 ms
64 bytes from 10.9.0.11: icmp_seq=3 ttl=64 time=0.050 ms
64 bytes from 10.9.0.11: icmp_seq=4 ttl=64 time=0.131 ms
64 bytes from 10.9.0.11: icmp_seq=5 ttl=64 time=0.098 ms
64 bytes from 10.9.0.11: icmp_seq=6 ttl=64 time=0.054 ms
64 bytes from 10.9.0.11: icmp_seq=7 ttl=64 time=0.050 ms
64 bytes from 10.9.0.11: icmp_seq=8 ttl=64 time=0.077 ms
```

```
root@1bc65e6feaa8:/# telnet 10.9.0.11
Trying 10.9.0.11...
^C
```

可知该现象的原因是路由器的过滤规则只允许icmp请求报文输入和icmp响应报文输入，ping的报文可以进行传输，而telnet的报文无法进行传输。

## Task 2.B: Protecting the Internal Network

清除table。

```
root@6c05b2e686e8:/# iptables -F
root@6c05b2e686e8:/# iptables -P OUTPUT ACCEPT
root@6c05b2e686e8:/# iptables -P INPUT ACCEPT
```

在router中输入以下命令。

```
root@6c05b2e686e8:/# iptables -A FORWARD -p icmp --icmp-type echo-request -d 10.9.0.5/24 -j ACCEPT
root@6c05b2e686e8:/# iptables -A FORWARD -p icmp --icmp-type echo-reply -d 192.168.60.0/24 -j ACCEPT
root@6c05b2e686e8:/# iptables -A FORWARD -p icmp --icmp-type echo-request -d 192.168.60/24 -j DROP
root@6c05b2e686e8:/# iptables -A INPUT -p icmp -j ACCEPT
root@6c05b2e686e8:/# iptables -A OUTPUT -p icmp -j ACCEPT
root@6c05b2e686e8:/# iptables -P FORWARD DROP
```

设置如下：

```
root@6c05b2e686e8:/# iptables -L
Chain INPUT (policy ACCEPT)
target     prot opt source               destination
ACCEPT    icmp -- anywhere            anywhere
Chain FORWARD (policy DROP)
target     prot opt source               destination
ACCEPT    icmp -- anywhere            10.9.0.0/24      icmp echo-request
ACCEPT    icmp -- anywhere            192.168.60.0/24   icmp echo-reply
DROP      icmp -- anywhere            192.168.60.0/24   icmp echo-request
Chain OUTPUT (policy ACCEPT)
target     prot opt source               destination
ACCEPT    icmp -- anywhere            anywhere
```

从外部主机 ping 路由器，可以 ping 通； ping 内部主机，不通。

```
root@1bc65e6feaa8:/# ping 10.9.0.11
PING 10.9.0.11 (10.9.0.11) 56(84) bytes of data.
64 bytes from 10.9.0.11: icmp_seq=1 ttl=64 time=0.168 ms
64 bytes from 10.9.0.11: icmp_seq=2 ttl=64 time=0.143 ms
64 bytes from 10.9.0.11: icmp_seq=3 ttl=64 time=0.262 ms
64 bytes from 10.9.0.11: icmp_seq=4 ttl=64 time=0.058 ms
64 bytes from 10.9.0.11: icmp_seq=5 ttl=64 time=0.090 ms
^C
--- 10.9.0.11 ping statistics ---
5 packets transmitted, 5 received, 0% packet loss, time 4097ms
rtt min/avg/max/mdev = 0.058/0.144/0.262/0.070 ms
root@1bc65e6feaa8:/# ping 192.168.60.5
PING 192.168.60.5 (192.168.60.5) 56(84) bytes of data.
^C
--- 192.168.60.5 ping statistics ---
8 packets transmitted, 0 received, 100% packet loss, time 7166ms
```

内部主机 ping 外部主机，可以 ping 通； telnet 外部主机，不通。

```
PING 10.9.0.5 (10.9.0.5) 56(84) bytes of data.
64 bytes from 10.9.0.5: icmp_seq=1 ttl=63 time=1.36 ms
64 bytes from 10.9.0.5: icmp_seq=2 ttl=63 time=0.121 ms
64 bytes from 10.9.0.5: icmp_seq=3 ttl=63 time=0.415 ms
64 bytes from 10.9.0.5: icmp_seq=4 ttl=63 time=0.083 ms
64 bytes from 10.9.0.5: icmp_seq=5 ttl=63 time=0.184 ms
64 bytes from 10.9.0.5: icmp_seq=6 ttl=63 time=0.503 ms
64 bytes from 10.9.0.5: icmp_seq=7 ttl=63 time=0.098 ms
64 bytes from 10.9.0.5: icmp_seq=8 ttl=63 time=0.080 ms
64 bytes from 10.9.0.5: icmp_seq=9 ttl=63 time=0.221 ms
64 bytes from 10.9.0.5: icmp_seq=10 ttl=63 time=0.675 ms
64 bytes from 10.9.0.5: icmp_seq=11 ttl=63 time=0.100 ms
64 bytes from 10.9.0.5: icmp_seq=12 ttl=63 time=0.091 ms
64 bytes from 10.9.0.5: icmp_seq=13 ttl=63 time=0.279 ms
64 bytes from 10.9.0.5: icmp_seq=14 ttl=63 time=0.069 ms
64 bytes from 10.9.0.5: icmp_seq=15 ttl=63 time=0.113 ms
64 bytes from 10.9.0.5: icmp_seq=16 ttl=63 time=0.144 ms
64 bytes from 10.9.0.5: icmp_seq=17 ttl=63 time=0.755 ms
64 bytes from 10.9.0.5: icmp_seq=18 ttl=63 time=0.406 ms
64 bytes from 10.9.0.5: icmp_seq=19 ttl=63 time=0.088 ms
^C
--- 10.9.0.5 ping statistics ---
19 packets transmitted, 19 received, 0% packet loss, time 18354ms
rtt min/avg/max/mdev = 0.069/0.304/1.363/0.322 ms
root@1d6463215ecb:/# telnet 10.9.0.5
Trying 10.9.0.5...
^C
```

## Task 2.C: Protecting Internal Servers

router中输入以下命令。

```
root@6c05b2e686e8:/# iptables -A FORWARD -p tcp --dport 23 -d 192.168.60.5 -j ACCEPT
root@6c05b2e686e8:/# iptables -A FORWARD -p tcp --sport 23 -s 192.168.60.5 -j ACCEPT
root@6c05b2e686e8:/# iptables -A FORWARD -d 10.9.0.0/24 -j DROP
root@6c05b2e686e8:/# iptables -A FORWARD -d 192.168.60.0/24 -j DROP
```

查看设置。

```
root@6c05b2e686e8:/# iptables -L
Chain INPUT (policy ACCEPT)
target     prot opt source               destination
Chain FORWARD (policy DROP)
target     prot opt source               destination
ACCEPT    tcp  --  anywhere             host1-192.168.60.5.net-192.168.60.0  tcp dpt:telnet
ACCEPT    tcp  --  host1-192.168.60.5.net-192.168.60.0  anywhere            tcp spt:telnet
DROP      all   --  anywhere             10.9.0.0/24
DROP      all   --  anywhere             192.168.60.0/24
Chain OUTPUT (policy ACCEPT)
target     prot opt source               destination
```

从外部主机(10.9.0.5)telnet 192.168.60.5，可以连接成功。

```
root@1bc65e6fea8:/# telnet 192.168.60.5
Trying 192.168.60.5...
Connected to 192.168.60.5.
Escape character is '^]'.
^Cc^C^C^C^CUbuntu 20.04.1 LTS
1d6463215ecb login: ^CConnection closed by foreign host.
root@1bc65e6fea8:/#
```

从外部主机(10.9.0.5)telnet 192.168.60.6，无法连接。

```
root@1bc65e6fea8:/# telnet 192.168.60.6
Trying 192.168.60.6...
^C
```

从内部主机(192.168.60.5)telnet 10.9.0.5，无法连接，内部主机(192.168.60.5)telnet 192.168.60.6，连接成功。

```
root@1d6463215ecb:/# telnet 10.9.0.5
Trying 10.9.0.5...
^C
root@1d6463215ecb:/# telnet 192.168.60.6
Trying 192.168.60.6...
Connected to 192.168.60.6.
Escape character is '^]'.
Ubuntu 20.04.1 LTS
8068735d0b54 login: ^CConnection closed by foreign host.
```

## Task 3: Connection Tracking and Stateful Firewall

### Task 3.A: Experiment with the Connection Tracking

docker重启，哈希值发生变化。

```
[07/26/21] seed@VM:~/.../Labsetup$ dockps
c4266f2dc484  host2-192.168.60.6
b17ef20a50a0  host3-192.168.60.7
9cdb090de1e5  hostA-10.9.0.5
6487420f2c41  host1-192.168.60.5
4ddde5e69374  seed-router
```

## ICMP experiment

在主机A (10.9.0.5) 上ping 192.168.60.5

```
PING 192.168.60.5 (192.168.60.5) 56(84) bytes of data.  
64 bytes from 192.168.60.5: icmp_seq=1 ttl=63 time=0.331 ms  
64 bytes from 192.168.60.5: icmp_seq=2 ttl=63 time=0.369 ms  
64 bytes from 192.168.60.5: icmp_seq=3 ttl=63 time=0.062 ms  
64 bytes from 192.168.60.5: icmp_seq=4 ttl=63 time=0.234 ms  
64 bytes from 192.168.60.5: icmp_seq=5 ttl=63 time=0.077 ms  
64 bytes from 192.168.60.5: icmp_seq=6 ttl=63 time=0.604 ms  
64 bytes from 192.168.60.5: icmp_seq=7 ttl=63 time=0.084 ms  
64 bytes from 192.168.60.5: icmp_seq=8 ttl=63 time=0.076 ms  
c^C  
--- 192.168.60.5 ping statistics ---  
8 packets transmitted, 8 received, 0% packet loss, time 7168ms  
rtt min/avg/max/mdev = 0.062/0.229/0.604/0.182 ms
```

查看连接状态， ICMP 的连接状态保持时间只有 30 秒左右。

```
root@4ddde5e69374:/# conntrack -L  
icmp 1 17 src=10.9.0.5 dst=192.168.60.5 type=8 code=0 id=29 src=192.168.60.  
5 dst=10.9.0.5 type=0 code=0 id=29 mark=0 use=1  
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
```

## UDP experiment

在主机 (192.168.60.5) 和主机A (10.9.0.5) 上建立UDP连接。

```
root@9cdb090de1e5:/# nc -u 192.168.60.5 9090  
hello  
■  
-----  
root@6487420f2c41:/# nc -lu 9090  
hello
```

查看连接状态， UDP 的连接状态保持时间和也只有 20~30 秒之间。

```
root@4ddde5e69374:/# conntrack -L  
udp 17 22 src=10.9.0.5 dst=192.168.60.5 sport=37135 dport=9090 [UNREPLIED]  
src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=37135 mark=0 use=1  
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
```

## TCP experiment

同上，建立TCP连接。

```
root@9cdb090de1e5:/# nc 192.168.60.5 9090  
hello  
■  
root@6487420f2c41:/# nc -l 9090  
hello  
■
```

TCP 的连接状态保持时间非常长。

```
root@4ddde5e69374:/# conntrack -L  
tcp 6 431993 ESTABLISHED src=10.9.0.5 dst=192.168.60.5 sport=44122 dport=9  
090 src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=44122 [ASSURED] mark=0 use=1  
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.  
root@4ddde5e69374:/# conntrack -L  
tcp 6 431988 ESTABLISHED src=10.9.0.5 dst=192.168.60.5 sport=44122 dport=9  
090 src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=44122 [ASSURED] mark=0 use=1  
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
```

## Task 3.B: Setting Up a Stateful Firewall

清理table，并输入如下命令。

```
root@4ddde5e69374:/# iptables -F
root@4ddde5e69374:/# iptables -A FORWARD -p tcp -m conntrack --ctstate ESTABLISHED,RELATED -j ACCEPT
root@4ddde5e69374:/# iptables -A FORWARD -p tcp --dport 23 -d 192.168.60.5 --syn -m conntrack --ctstate NEW -j ACCEPT
root@4ddde5e69374:/# iptables -A FORWARD -p tcp --dport 23 -d 10.9.0.0/24 --syn -m conntrack --ctstate NEW -j ACCEPT
root@4ddde5e69374:/# iptables -P FORWARD DROP
```

从外部主机(10.9.0.5)telnet 192.168.60.5。, 连接成功。

```
root@9cdb090de1e5:/# telnet 192.168.60.5
Trying 192.168.60.5...
Connected to 192.168.60.5.
Escape character is '^]'.
Ubuntu 20.04.1 LTS
6487420f2c41 login: ^CConnection closed by foreign host.
```

telnet 192.168.60.6。, 连接失败。

```
root@9cdb090de1e5:/# telnet 192.168.60.6
Trying 192.168.60.6...
^C
```

从内部主机(192.168.60.5)telnet 10.9.0.5 和 192.168.60.6，连接成功。

```
root@6487420f2c41:/# telnet 10.9.0.5
Trying 10.9.0.5...
Connected to 10.9.0.5.
Escape character is '^]'.
Ubuntu 20.04.1 LTS
9cdb090de1e5 login: c^CConnection closed by foreign host.
root@6487420f2c41:/# telnet 192.168.60.6
Trying 192.168.60.6...
Connected to 192.168.60.6.
Escape character is '^]'.
Ubuntu 20.04.1 LTS
c4266f2dc484 login: ^CConnection closed by foreign host.
```

不利用连接跟踪机制的过滤规则仅对数据包的首部进行检查，其优点是处理速度快，缺点是无法定义精细的规则、不适合复杂的访问机制；而利用连接跟踪机制的过滤规则对数据包的状态也进行检查，其优点是能够定义更加严格的规则、应用范围更广、安全性更高，缺点是无法对数据包的内容进行识别。

## Task 4: Limiting Network Traffic

在router上利用iptables命令，创建流量限制规则如下。

```
[07/27/21]seed@VM:~/.../Labsetup$ docksh 4d
root@4ddde5e69374:/# iptables -A FORWARD -s 10.9.0.5 -m limit --limit 10/minute
--limit-burst 5 -j ACCEPT
root@4ddde5e69374:/# iptables -A FORWARD -s 10.9.0.5 -j DROP
```

从外部(10.9.0.5)ping 192.168.60.5，得到结果如下，可知能够连接，可以观察到前六个包的速度很快，后面发包速度变慢。

```
root@9cdb090de1e5:/# ping 192.168.60.5
PING 192.168.60.5 (192.168.60.5) 56(84) bytes of data.
64 bytes from 192.168.60.5: icmp_seq=1 ttl=63 time=0.153 ms
64 bytes from 192.168.60.5: icmp_seq=2 ttl=63 time=0.380 ms
64 bytes from 192.168.60.5: icmp_seq=3 ttl=63 time=0.098 ms
64 bytes from 192.168.60.5: icmp_seq=4 ttl=63 time=0.253 ms
64 bytes from 192.168.60.5: icmp_seq=5 ttl=63 time=0.262 ms
64 bytes from 192.168.60.5: icmp_seq=7 ttl=63 time=0.084 ms
64 bytes from 192.168.60.5: icmp_seq=13 ttl=63 time=0.310 ms
64 bytes from 192.168.60.5: icmp_seq=19 ttl=63 time=0.225 ms
64 bytes from 192.168.60.5: icmp_seq=25 ttl=63 time=0.084 ms
64 bytes from 192.168.60.5: icmp_seq=31 ttl=63 time=0.106 ms
c^C
--- 192.168.60.5 ping statistics ---
31 packets transmitted, 10 received, 67.7419% packet loss, time 30724ms
rtt min/avg/max/mdev = 0.084/0.195/0.380/0.099 ms
```

如果只执行第一条命令，从外部(10.9.0.5)ping 192.168.60.5，可以观察到和平时的发包速度一样，因为iptables默认的 FORWARD 表是接受所有包，所以如果不写第二条命令，发包会正常进行。

```
root@4ddde5e69374:/# iptables -A FORWARD -s 10.9.0.5 -m limit --limit 10/minute
--limit-burst 5 -j ACCEPT
root@9cdb090de1e5:/# ping 192.168.60.5
PING 192.168.60.5 (192.168.60.5) 56(84) bytes of data.
64 bytes from 192.168.60.5: icmp_seq=1 ttl=63 time=0.112 ms
64 bytes from 192.168.60.5: icmp_seq=2 ttl=63 time=0.084 ms
64 bytes from 192.168.60.5: icmp_seq=3 ttl=63 time=0.252 ms
64 bytes from 192.168.60.5: icmp_seq=4 ttl=63 time=0.069 ms
64 bytes from 192.168.60.5: icmp_seq=5 ttl=63 time=0.084 ms
64 bytes from 192.168.60.5: icmp_seq=6 ttl=63 time=0.236 ms
64 bytes from 192.168.60.5: icmp_seq=7 ttl=63 time=0.520 ms
64 bytes from 192.168.60.5: icmp_seq=8 ttl=63 time=0.105 ms
64 bytes from 192.168.60.5: icmp_seq=9 ttl=63 time=0.090 ms
64 bytes from 192.168.60.5: icmp_seq=10 ttl=63 time=0.345 ms
64 bytes from 192.168.60.5: icmp_seq=11 ttl=63 time=0.082 ms
64 bytes from 192.168.60.5: icmp_seq=12 ttl=63 time=0.089 ms
^C
--- 192.168.60.5 ping statistics ---
12 packets transmitted, 12 received, 0% packet loss, time 11274ms
rtt min/avg/max/mdev = 0.069/0.172/0.520/0.134 ms
```

## Task 5: Load Balancing

### nth mode

在router上利用iptables命令，采用nth模式创建负载均衡规则如下：

```
root@4ddde5e69374:/# iptables -t nat -A PREROUTING -p udp --dport 8080 -m statistic --mode nth --every 3 --packet 0 -j DNAT --to-destination 192.168.60.5:8080
root@4ddde5e69374:/# iptables -t nat -A PREROUTING -p udp --dport 8080 -m statistic --mode nth --every 3 --packet 1 -j DNAT --to-destination 192.168.60.6:8080
root@4ddde5e69374:/# iptables -t nat -A PREROUTING -p udp --dport 8080 -m statistic --mode nth --every 3 --packet 2 -j DNAT --to-destination 192.168.60.7:8080
```

发包情况如下，按顺序 hello\_1 被发送到 192.168.60.5 8080，hello\_2 被发送到 192.168.60.6 8080，hello\_3 被发送到 192.168.60.7 8080。

```
root@9cdb090de1e5:/# echo hello1 | nc -u 10.9.0.11 8080
^C
root@9cdb090de1e5:/# echo hello2 | nc -u 10.9.0.11 8080
root@9cdb090de1e5:/# echo hello2 | nc -u 10.9.0.11 8080
^C
root@9cdb090de1e5:/# echo hello3 | nc -u 10.9.0.11 8080
^C
root@c4266f2dc484:/# nc -luk 8080
hello1
-----
root@b17ef20a50a0:/# nc -luk 8080
hello2
-----
root@6487420f2c41:/# nc -luk 8080
hello3
```

### random mode

清除之前的iptables规则，router中输入以下规则，等概率发送数据。

```
root@4ddde5e69374:/# iptables -t nat -A PREROUTING -p udp --dport 8080 -m statistic --mode random --probability 0.33 -j DNAT --to-destination 192.168.60.5:8080
root@4ddde5e69374:/# iptables -t nat -A PREROUTING -p udp --dport 8080 -m statistic --mode random --probability 0.33 -j DNAT --to-destination 192.168.60.6:8080
root@4ddde5e69374:/# iptables -t nat -A PREROUTING -p udp --dport 8080 -m statistic --mode random --probability 0.33 -j DNAT --to-destination 192.168.60.7:8080
```

虽然是等概率发送数据，但每个主机收到的数量各不相同，甚至有的差异较大，当样本数量足够多时，应该是趋于平均的。

```
root@6487420f2c41:/# nc -luk 8080
hello

root@c4266f2dc484:/# nc -luk 8080
hello
hello
hello

root@b17ef20a50a0:/# nc -luk 8080
hello
hello
hello
```